

Foundations of Deep Learning, Winter Term 2021/22

Week 12: Variational Autoencoders and Generative Adversarial Networks

## Variational Autoencoders and Generative Adversarial Networks

Frank Hutter    Abhinav Valada

University of Freiburg



# Overview of Week 12

- 1 Introduction to Autoencoders
- 2 Sparse and Denoising Autoencoders
- 3 Generative Models and the Variational Autoencoder
- 4 Generative Adversarial Networks
- 5 Summary, Further Reading, References

# Lecture Overview

- 1 Introduction to Autoencoders
- 2 Sparse and Denoising Autoencoders
- 3 Generative Models and the Variational Autoencoder
- 4 Generative Adversarial Networks
- 5 Summary, Further Reading, References

Foundations of Deep Learning, Winter Term 2021/22

Week 12: Variational Autoencoders and Generative Adversarial Networks

## Introduction to Autoencoders

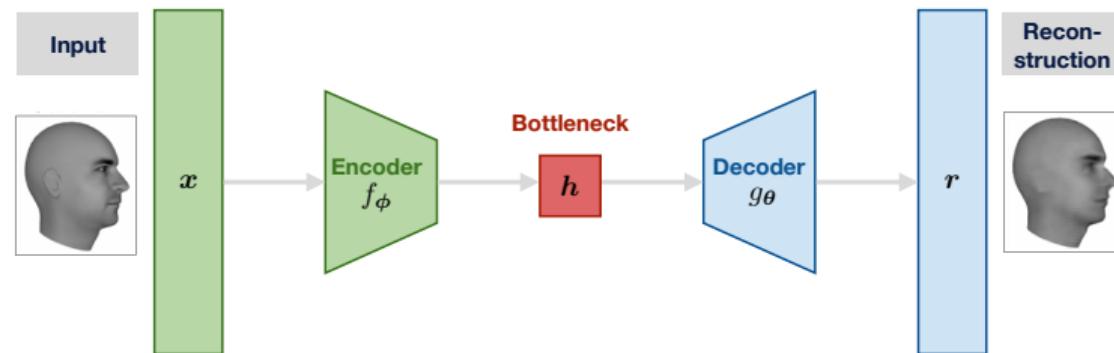
Frank Hutter Abhinav Valada

University of Freiburg



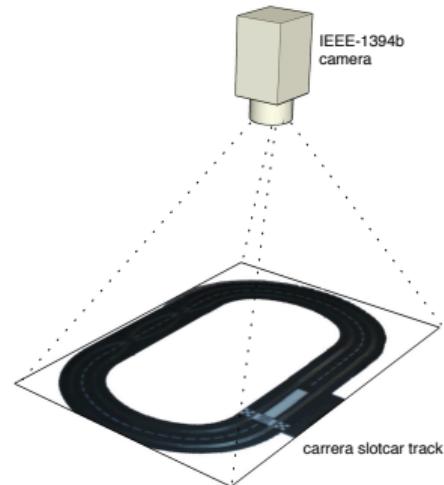
# Autoencoders

- Neural networks trained to reconstruct their input.



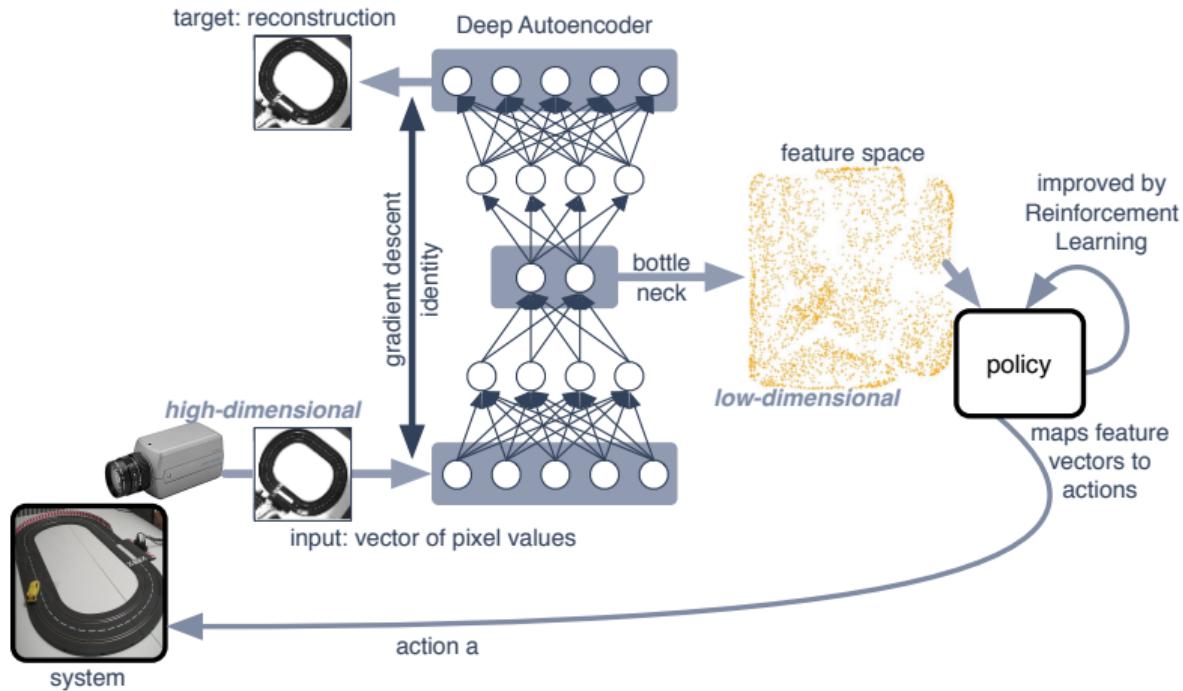
- **Architecture** and/or **regularization** is chosen to **prevent** simple copying. → Forces the learning process to **prioritize** what information to keep and what to throw out.
- Hidden layer  $h$  describes a compressed **code** to represent the input.
- Two parts: **encoder** function  $h = f_\phi(x)$ , **decoder** function  $r = g_\theta(h)$
- Learning process: minimize  $L(x, g_\theta(f_\phi(x)))$

# Example Application: Learning to Control a Slotcar from Pixel Input



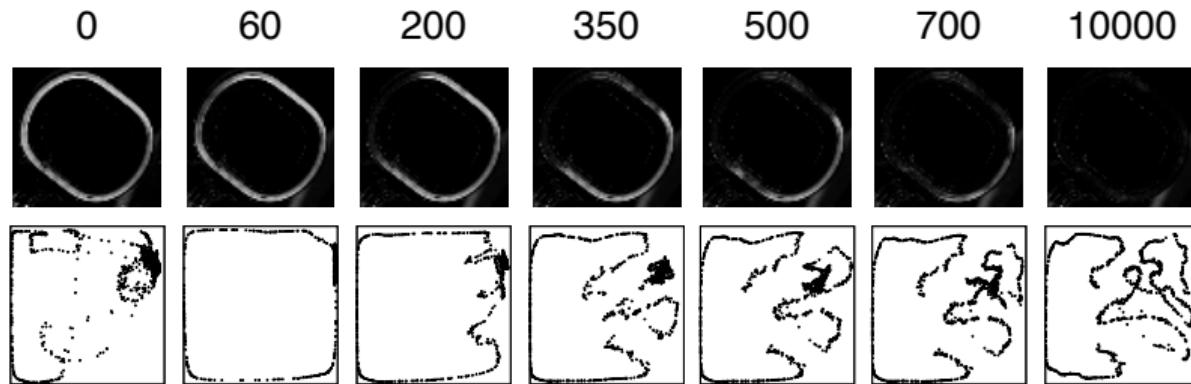
[Image source: Lange et al., 2012]

# Overall Setup



[Image source: Lange et al., 2012]

# Evolution of the Latent Space



[Image source: Lange et al., 2012]

## Questions to Answer for Yourself / Discuss with Friends

- Repetition: How do you train an autoencoder and why is it useful?
- Application of what you just learned: Can you think of other applications where using an autoencoder would make sense?

# Lecture Overview

- 1 Introduction to Autoencoders
- 2 Sparse and Denoising Autoencoders
- 3 Generative Models and the Variational Autoencoder
- 4 Generative Adversarial Networks
- 5 Summary, Further Reading, References

Foundations of Deep Learning, Winter Term 2021/22

Week 12: Variational Autoencoders and Generative Adversarial Networks

## Sparse and Denoising Autoencoders

Frank Hutter    Abhinav Valada

University of Freiburg

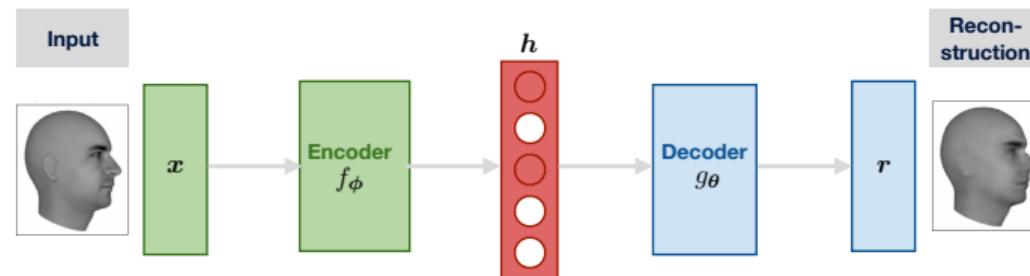


# Regularized Autoencoders

- Allow for dimension of  $h \geq$  dimension of  $x$ . → **overcomplete** representation
- Model capacity can be chosen based on the complexity of the data distribution.
- Prevent simple copying of the input (and therefore pressure to learn sth. useful about the input data), e.g., by:
  - adding **sparsity** penalty on the hidden layer activations.
  - changing the reconstruction cost to a **denoising** objective.
  - penalizing derivatives of hidden layer activations w.r.t input.  
→ **contractive autoencoder** (see chapter 14.7 [Goodfellow et al., 2016] for more details)

# Sparse Autoencoders

- Add a **sparsity penalty**  $\Omega(\mathbf{h})$  on the code  $\mathbf{h}$  in the hidden layer to the loss function:  
 $L(\mathbf{x}, g_\theta(f_\phi(\mathbf{x}))) + \Omega(\mathbf{h})$
- Often used as an unsupervised pre-processing step for downstream supervised processing, e.g., classification.
- Sparsity as regularization to avoid overfitting and increase generalization.
- Can be interpreted as **approximating maximum likelihood training** of a **generative model** with visible variables  $\mathbf{x}$  and latent variables  $\mathbf{h}$ .
- Joint distribution:  $p_{model}(\mathbf{x}, \mathbf{h}) = p_{model}(\mathbf{h})p_{model}(\mathbf{x}|\mathbf{h})$

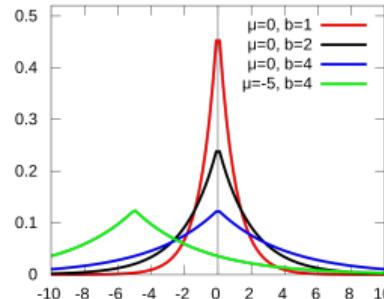


## Sparse Autoencoders (cont.)

- Log likelihood can be decomposed as:

$$\log p_{model}(\mathbf{x}) = \log \sum_{\mathbf{h}} p_{model}(\mathbf{h}, \mathbf{x})$$

- Sparse autoencoder will approximate this sum with a point estimate of one highly likely value for  $\mathbf{h}$ .
- We maximize:  $\log p_{model}(\mathbf{h}, \mathbf{x}) = \log p_{model}(\mathbf{h}) + \log p_{model}(\mathbf{x}|\mathbf{h})$
- Choosing a Laplace prior implements sparsity:  $p_{model}(h_i) = \frac{\lambda}{2} \exp^{-\lambda|h_i|}$



[Image source: [https://commons.wikimedia.org/wiki/File:Laplace\\_pdf\\_mod.svg](https://commons.wikimedia.org/wiki/File:Laplace_pdf_mod.svg)]

## Sparse Autoencoders (cont.)

The (negative) log-prior of

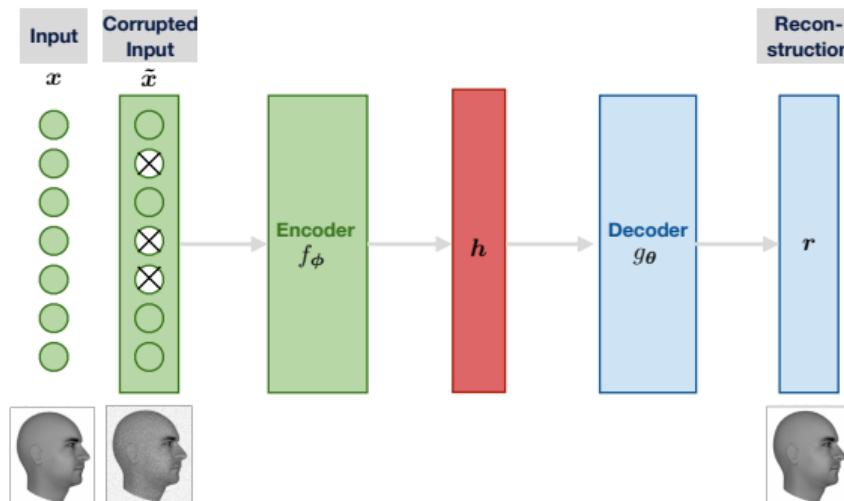
$$p_{model}(h_i) = \frac{\lambda}{2} \exp^{-\lambda|h_i|}$$

corresponds to an absolute value penalty  $\Omega(\mathbf{h}) = \lambda \sum_i |h_i|$ :

$$-\log p_{model}(\mathbf{h}) = \sum_i (\lambda|h_i| - \log \frac{\lambda}{2}) = \Omega(\mathbf{h}) + const$$

# Denoising Autoencoders

- Rather than adding an explicit regularization term, denoising autoencoders (DAE) learn to **undo the effect of noise corruption** when learning to reconstruct the input.
- The DAE minimizes the loss function  $L(x, g_\theta(f_\phi(\tilde{x})))$  with  $\tilde{x}$  being a corrupted version of input  $x$ .
- To implement this, we introduce a corruption process  $C(\tilde{x}|x)$ .

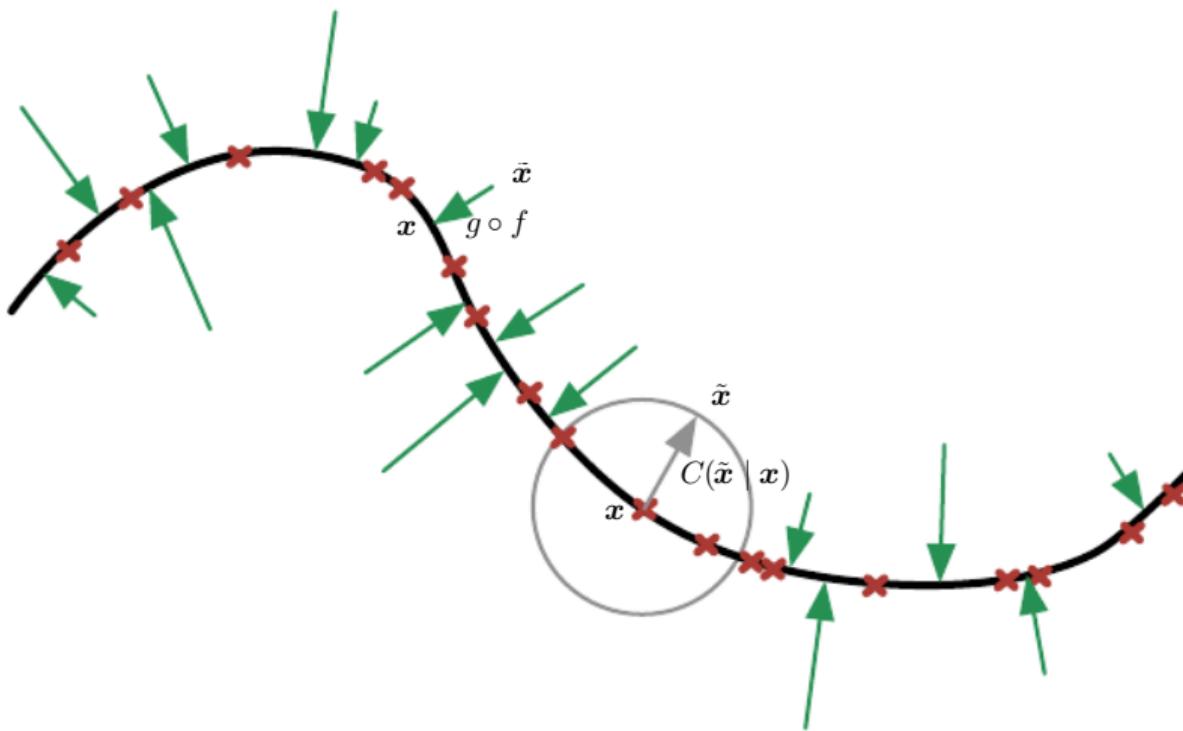


## Denoising Autoencoders (cont.)

- The DAE learns a **reconstruction distribution**  $p_{\text{reconstruct}}(\mathbf{x}|\tilde{\mathbf{x}})$  from training pairs  $(\mathbf{x}, \tilde{\mathbf{x}})$  as follows:
  - ➊ Sample a training example  $\mathbf{x}$  from the training data.
  - ➋ Sample a corrupted version  $\tilde{\mathbf{x}}$  from  $C(\tilde{\mathbf{x}}|\mathbf{x} = \mathbf{x})$ .
  - ➌ Use  $(\mathbf{x}, \tilde{\mathbf{x}})$  as a training example for estimating the autoencoder reconstruction distribution  $p_{\text{reconstruct}}(\mathbf{x}|\tilde{\mathbf{x}}) = p_{\text{decoder}}(\mathbf{x}|\mathbf{h})$  with  $\mathbf{h}$  the output of encoder  $f_\phi(\tilde{\mathbf{x}})$  and  $p_{\text{decoder}}$  typically defined by a decoder  $g_\theta(\mathbf{h})$ .
- Perform gradient-based approximate minimization on the negative log-likelihood  $-\log p_{\text{decoder}}(\mathbf{x}|\mathbf{h})$  (same techniques as for other types of FFN, as long as encoder is deterministic).
- The DAE therefore performs stochastic gradient descent on:

$$-\mathbb{E}_{\mathbf{x} \sim \hat{p}_{\text{data}}(\mathbf{x})} \mathbb{E}_{\tilde{\mathbf{x}} \sim C(\tilde{\mathbf{x}}|\mathbf{x})} \log p_{\text{decoder}}(\mathbf{x}|\mathbf{h} = f(\tilde{\mathbf{x}}))$$

## Denoising Autoencoders (cont.)



[Image source: Goodfellow et al., 2016, p. 512, Fig. 14.4]

## Questions to Answer for Yourself / Discuss with Friends

- Repetition: Why do we need regularization in overcomplete autoencoders?
- Repetition: What is the difference between a sparse and a denoising autoencoder?

# Lecture Overview

- 1 Introduction to Autoencoders
- 2 Sparse and Denoising Autoencoders
- 3 Generative Models and the Variational Autoencoder
- 4 Generative Adversarial Networks
- 5 Summary, Further Reading, References

Foundations of Deep Learning, Winter Term 2021/22

Week 12: Variational Autoencoders and Generative Adversarial Networks

## Generative Models and the Variational Autoencoder

Frank Hutter    Abhinav Valada

University of Freiburg



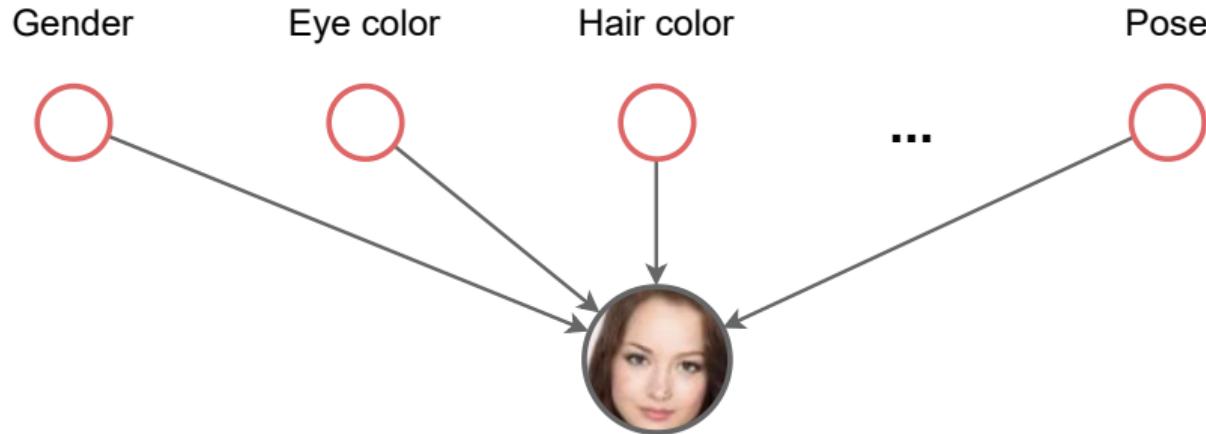
## Illustration: Factors of Variation in Images



Factors that affect the **variability** in images of faces  $x$  are, e.g., gender, eye color, hair color, pose, etc. All we can observe in the image are the pixels though, the factors of variation are not explicitly available, i.e., they are **latent**.

[Image source: Berthelot et al., 2017]

## Illustration: Factors of Variation in Images (cont.)



- Idea: explicitly model these factors using latent variables  $z$
- If  $z$  reflects the causal factors that determine the appearance of the pixels in  $x$ , then  $p(x|z)$  should be of an easier form of trying to model all the possible sources of variation in  $p(x)$  directly.

[Image (face) source: Berthelot et al., 2017]

# Latent Variable Generative Models

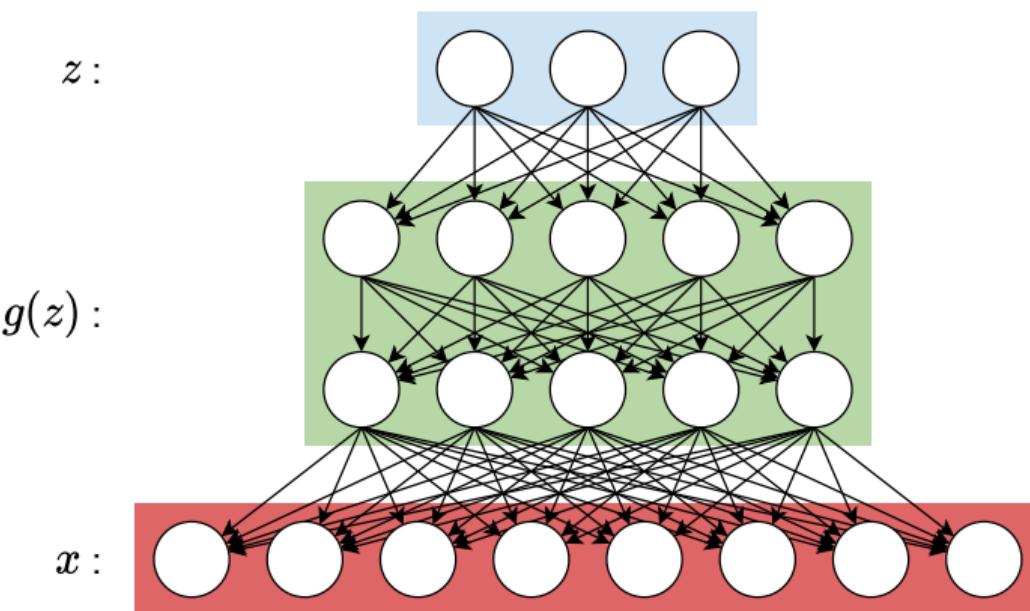
- Assuming set of **visible** variables  $\mathbf{x}$  and **latent** variables  $\mathbf{z}$  our goal is to model the distribution of the training data :

$$p(\mathbf{x}) = \int p(\mathbf{x}, \mathbf{z}) d\mathbf{z} = \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z}) d\mathbf{z}$$

- This breaks down the modeling problem in a hierarchical way: a context / explanatory factors  $p(\mathbf{z})$  of our data, and a generating distribution  $p(\mathbf{x}|\mathbf{z})$  which is conditioned on them.
- This is used in a variety of probabilistic models, e.g., **Gaussian mixture models**, **topic models**, etc.
- Once we have a good approximation of  $p(\mathbf{x})$ , we can **sample** from it and generate new data that is **very similar, but different** from our original training data.

# Variational Autoencoder (VAE) Approach

Use neural networks as **deterministic parts** within a graphical model to be able to express complex relationship between latent and visible variables.



# The Problem: How to Do Inference?

- How do we get  $z$  values that are related to our input data?
- This is the problem of inferring the **posterior** distribution  $p(z|x)$ .
- Bayes formula says:

$$p(z|x) = \frac{p(x|z)p(z)}{p(x)} = \frac{p(x|z)p(z)}{\int p(x,z)dz}$$

- Unfortunately, the integral in the denominator is intractable!  
→ Have to resort to **approximate inference** methods

# Variational Inference in the VAE

- Idea: Approximate the posterior with a simpler distribution  $q(\mathbf{z}|\mathbf{x})$  with its own set of parameters and optimize them to get as close as possible to  $p(\mathbf{z}|\mathbf{x})$ .
- The difference between the approximate and the real posterior can be measured by the KL-divergence:

$$\begin{aligned} KL(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) &= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \log \frac{q(\mathbf{z}|\mathbf{x})}{p(\mathbf{z}|\mathbf{x})} \right] \\ &= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} \left[ \log q(\mathbf{z}|\mathbf{x}) - \log \frac{p(\mathbf{z}, \mathbf{x})}{p(\mathbf{x})} \right] \\ &= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} [\log q(\mathbf{z}|\mathbf{x}) - \log p(\mathbf{z}, \mathbf{x})] + \log p(\mathbf{x}) \end{aligned}$$

- We cannot compute this KL-divergence directly because of the intractable integration needed to get  $p(\mathbf{x})$ . We can, however, manipulate  $q$ .

## Variational Inference in the VAE (cont.)

- Observe that

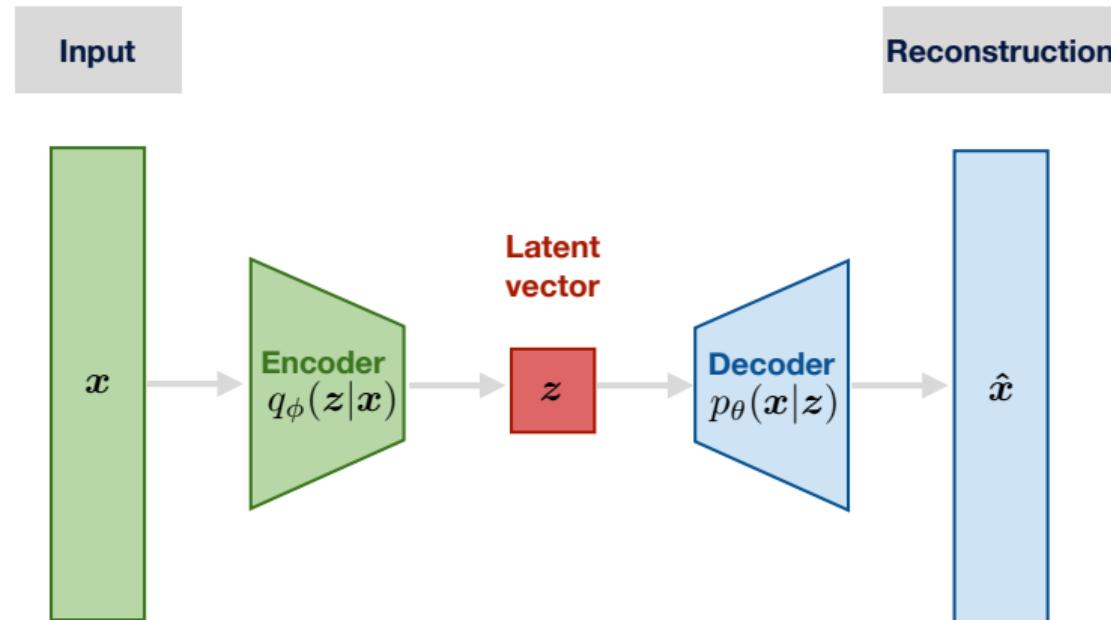
$$\log p(\mathbf{x}) = -\mathbb{E}_{q(\mathbf{z}|\mathbf{x})} [\log q(\mathbf{z}|\mathbf{x}) - \log p(\mathbf{z}, \mathbf{x})] + KL(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x}))$$

and that the KL-divergence has a minimum of zero.

- This means that  $\mathcal{L}(\mathbf{x}, q) = \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} [\log p(\mathbf{z}, \mathbf{x}) - \log q(\mathbf{z}|\mathbf{x})]$  is a lower bound on  $\log p(\mathbf{x})$ .
- It is called the **evidence lower bound** (ELBO) or the **variational lower bound**.
- For an appropriate choice of  $q$ ,  $\mathcal{L}$  is tractable to compute. We maximize  $\mathcal{L}$  to make the bound as tight as possible (KL-divergence will be zero at the maximum).

# Variational Autoencoder

The **variational autoencoder** (VAE) uses two neural networks: an **encoder** network with parameters  $\phi$  that learns to approximate the posterior, and a **decoder** network with parameters  $\theta$  that learns to reconstruct the input.



## Variational Autoencoder (cont.)

- With this parametrization, the ELBO for variational inference is:

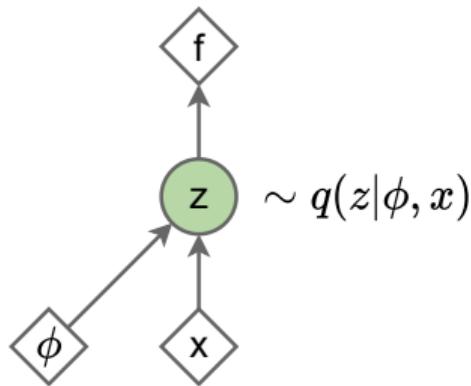
$$\begin{aligned}\mathcal{L}(\theta, \phi, \mathbf{x}) &= \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{z}, \mathbf{x}) - \log q_\phi(\mathbf{z} | \mathbf{x})] \\ &= \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{x}|\mathbf{z}) + \log p_\theta(\mathbf{z}) - \log q_\phi(\mathbf{z} | \mathbf{x})] \\ &= \underbrace{-KL(\log q_\phi(\mathbf{z}|\mathbf{x}) || \log p_\theta(\mathbf{z}))}_{\text{regularization term}} + \underbrace{\mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{x}|\mathbf{z})]}_{\text{reconstruction term}}\end{aligned}$$

- We approximate the expectation in the ELBO by sampling from posterior over  $\mathbf{z}$ , but sampling is not a continuous operation – how can we backprop through this step?
- Idea:** Reparametrize to express our random variable  $\mathbf{z}$  as a deterministic variable and an external, static noise source expressed in random variable  $\epsilon$ .

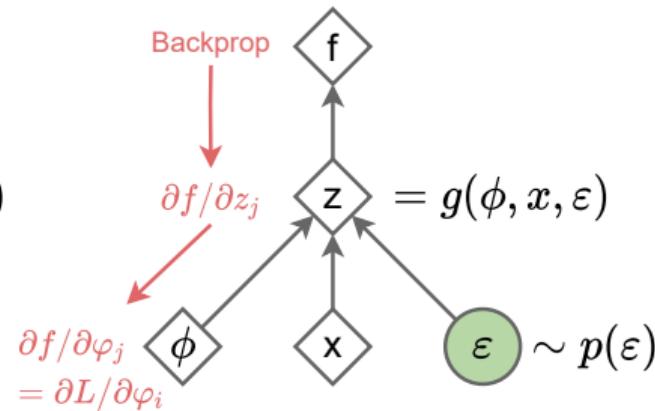
# Reparametrization Trick

Let  $z \in \mathbb{R}$  and  $q_\phi(z|x) = \mathcal{N}(z; \mu_z(x), \sigma_z(x)^2 \odot \mathbf{I})$

Original form:



Reparameterized form:

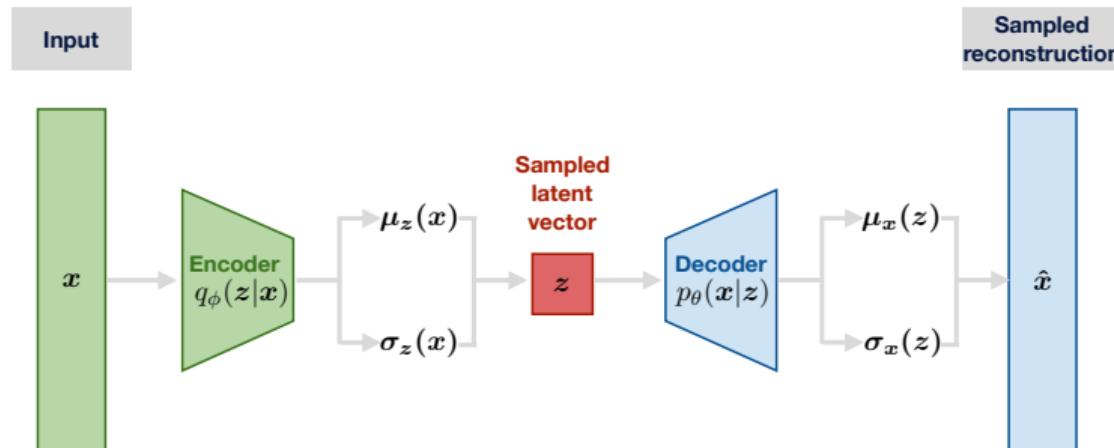


- Use parametrization  $z = \mu_z(x) + \sigma_z(x) \odot \epsilon_z$  where  $\epsilon_z \sim \mathcal{N}(0, \mathbf{I})$
- Optional: use parametrization  $x = \mu_x(z) + \sigma_x(z) \odot \epsilon_x$  where  $\epsilon_x \sim \mathcal{N}(0, \mathbf{I})$

# Training with Backpropagation

Now, with this **reparametrization**, both the **generative** model  $p_\theta(\mathbf{x}|\mathbf{z})$  and the **inference** model  $q_\phi(\mathbf{z}|\mathbf{x})$  can be trained together using **backpropagation** to optimize the ELBO as objective function:

$$\mathcal{L}(\theta, \phi, \mathbf{x}) = -KL(\log q_\phi(\mathbf{z}|\mathbf{x}) || \log p_\theta(\mathbf{z})) + \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{x}|\mathbf{z})]$$



# VAE Training Pseudocode

---

**Algorithm 1** Minibatch version of the Auto-Encoding VB (AEVB) algorithm. Either of the two SGVB estimators in section 2.3 can be used. We use settings  $M = 100$  and  $L = 1$  in experiments.

---

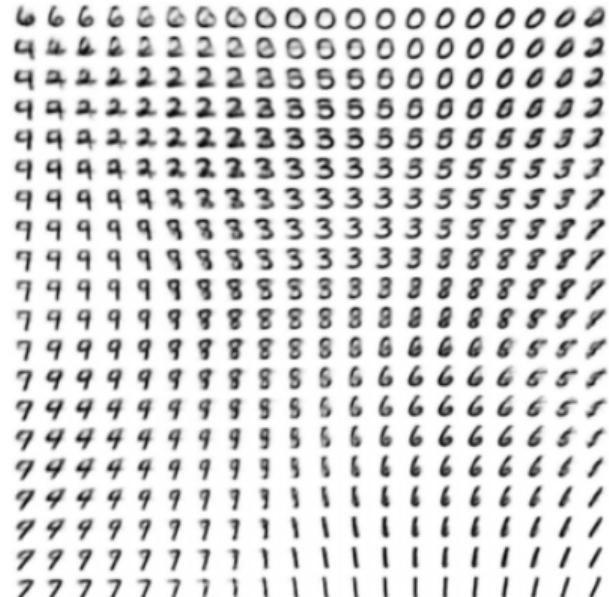
```
 $\theta, \phi \leftarrow$  Initialize parameters  
repeat  
   $\mathbf{X}^M \leftarrow$  Random minibatch of  $M$  datapoints (drawn from full dataset)  
   $\epsilon \leftarrow$  Random samples from noise distribution  $p(\epsilon)$   
   $\mathbf{g} \leftarrow \nabla_{\theta, \phi} \tilde{\mathcal{L}}^M(\theta, \phi; \mathbf{X}^M, \epsilon)$  (Gradients of minibatch estimator (8))  
   $\theta, \phi \leftarrow$  Update parameters using gradients  $\mathbf{g}$  (e.g. SGD or Adagrad [DHS10])  
until convergence of parameters  $(\theta, \phi)$   
return  $\theta, \phi$ 
```

---

[Algorithm source: Kingma and Welling, 2013]

# Learning Smooth Latent Spaces

Training a VAE with a two-dimensional latent code on MNIST (left) and the Frey Face (right) dataset, we get the following two latent spaces by varying  $z$  on a uniform grid.



[Image source: Kingma and Welling, 2013]

## Questions to Answer for Yourself / Discuss with Friends

- Repetition: What is the idea behind generative models?
- Repetition: What is the reparametrization trick?
- Repetition: How do you train a VAE?

# Lecture Overview

- 1 Introduction to Autoencoders
- 2 Sparse and Denoising Autoencoders
- 3 Generative Models and the Variational Autoencoder
- 4 Generative Adversarial Networks
- 5 Summary, Further Reading, References

Foundations of Deep Learning, Winter Term 2021/22

Week 12: Variational Autoencoders and Generative Adversarial Networks

## Generative Adversarial Networks

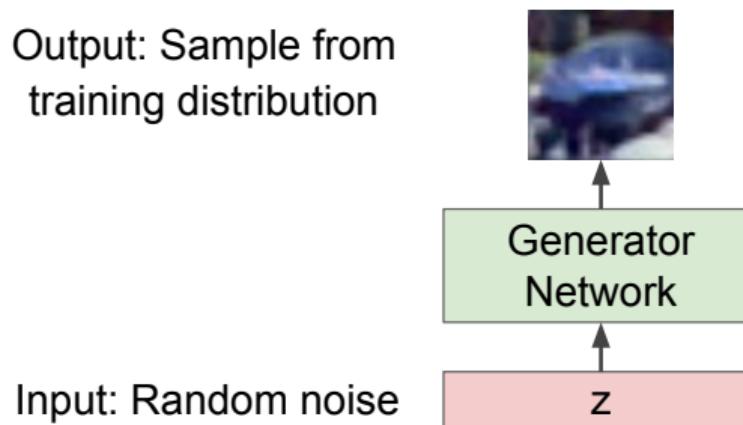
Frank Hutter    Abhinav Valada

University of Freiburg



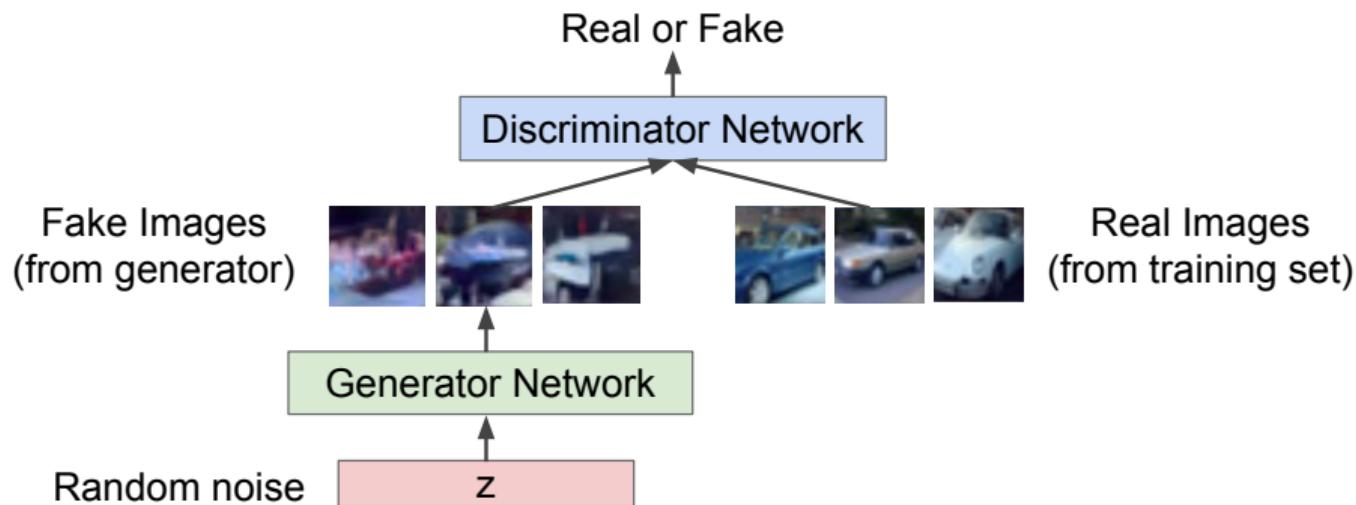
# Generative Adversarial Networks

- **Motivation:** GANs are generative models, similar to VAEs but drop the assumption of optimizing the lower variational bound. Hence they often produce better results.
- Want to sample from complex, high-dimensional training distribution. No direct way to do this!
- Sample from a simple distribution, e.g., random noise. Learn transformation to training distribution.



# Training GANs: Two-Player Game

- Generator network: try to fool the discriminator by generating real-looking images.
- Discriminator network: try to distinguish between real and fake images.



## Training GANs: Two-Player Game (cont.)

- Train jointly in minimax game with the objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

- Discriminator outputs likelihood in (0,1) of real image.
- Discriminator ( $\theta_d$ ) wants to maximize objective such that  $D(x)$  is close to 1 (real) and  $D(G(z))$  is close to 0 (fake).
- Generator ( $\theta_g$ ) wants to minimize objective such that  $D(G(z))$  is close to 1 (discriminator is fooled into thinking generated  $G(z)$  is real).

# Training GANs: Two-Player Game (cont.)

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Alternate between:

1. Gradient ascent on discriminator

$$\max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

2. Gradient descent on generator

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$

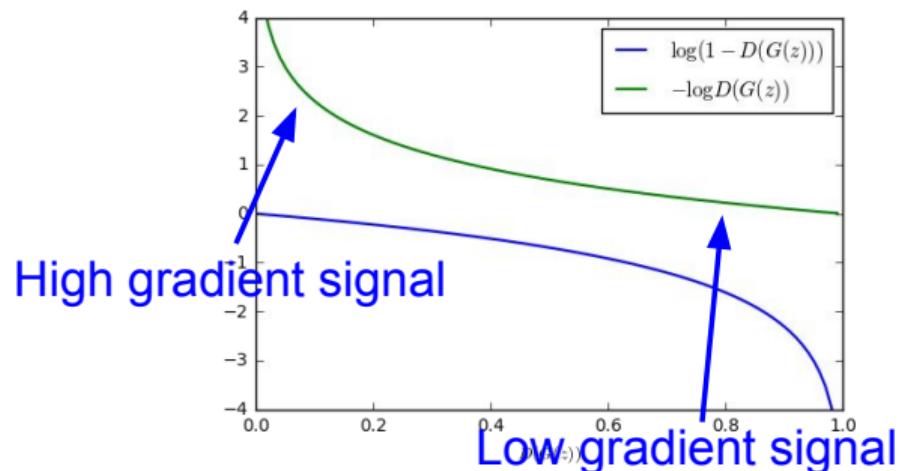
# A Better Cost Function

Original minmax cost:

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$

Modified generator cost:

$$\max_{\theta_d} \mathbb{E}_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))$$



# GAN Training Algorithm

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D_{\theta_d}(\mathbf{x}^{(i)}) + \log(1 - D_{\theta_d}(G_{\theta_g}(\mathbf{z}^{(i)}))) \right]$$

**end for**

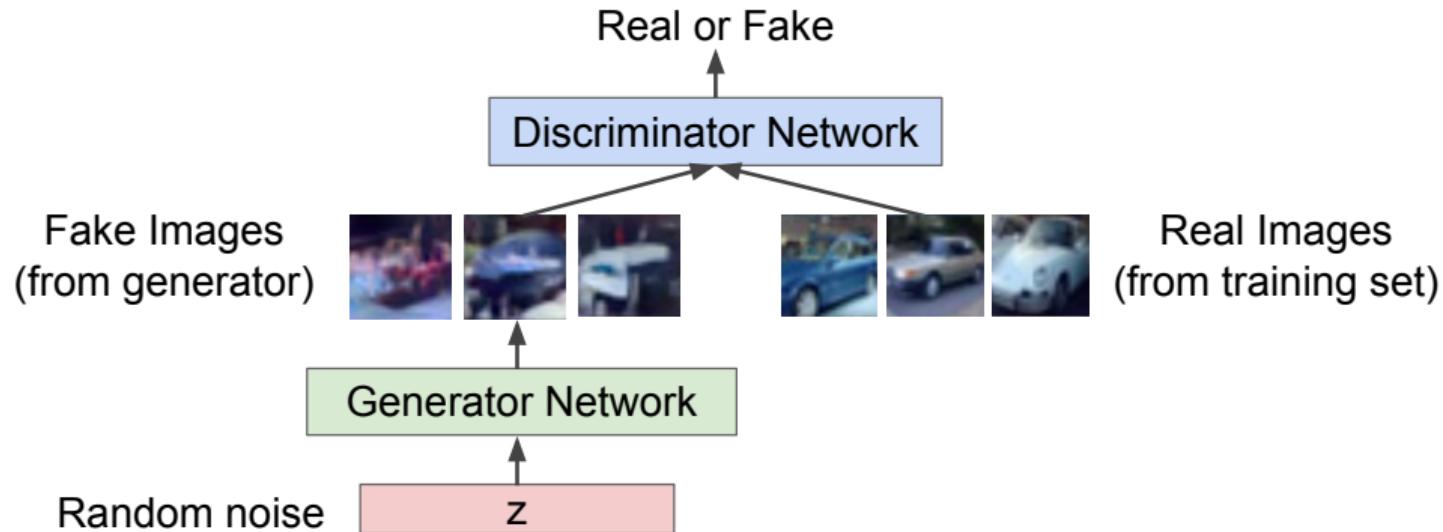
- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Update the generator by ascending its stochastic gradient (improved objective):

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log(D_{\theta_d}(G_{\theta_g}(\mathbf{z}^{(i)})))$$

**end for**

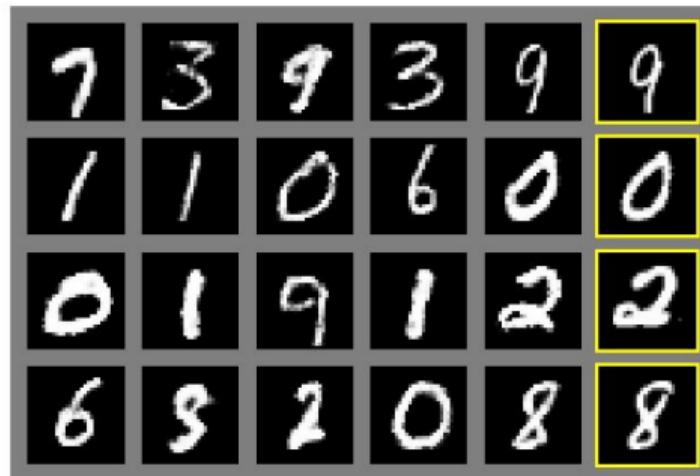
# Training GANs: Two-Player Game (cont.)

- Generator network: try to fool the discriminator by generating real-looking images.
- Discriminator network: try to distinguish between real and fake images.
- After training, use generator network to generate new images.



# Generative Adversarial Nets

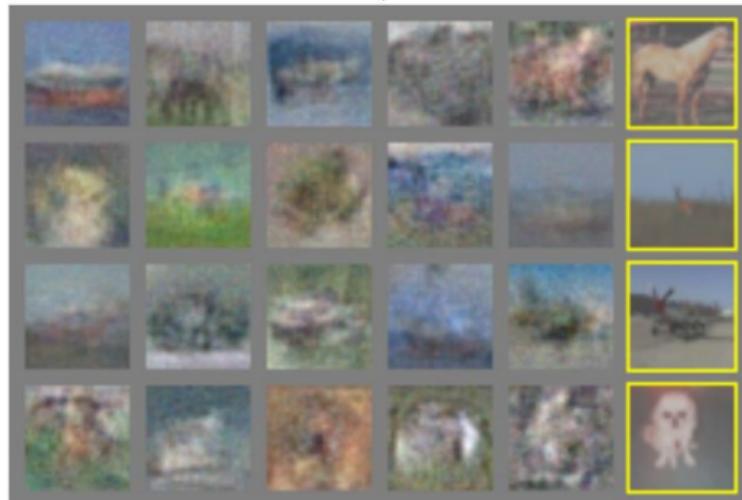
Generated Samples



[Image source: Goodfellow et al., 2014]

## Generative Adversarial Nets (cont.)

Generated Samples (CIFAR-10)



[Image source: Goodfellow et al., 2014]

# Generative Adversarial Nets – Convolutional Architectures

Architecture guidelines for stable deep convolutional GANs:

- Generator is an upsampling network with fractionally-strided convolutions.
- Discriminator is a convolutional network.
- Replace any pooling layers with strided convolutions (discriminator) and fractional-strided convolutions (generator).
- Use batchnorm in both the generator and discriminator.
- Remove fully connected hidden layers for deeper architectures.
- Use ReLU activation in generator for all layers except for the output, which uses tanh.
- Use LeakyReLU activation in the discriminator for all layers.

[Items 3 to 7 are from: Radford et al., 2016]

# Generative Adversarial Nets – Convolutional Architectures (cont.)



[Image source: Radford et al., 2016]

# 2017: Explosion of GANs

## Better training and generation



(a) Church outdoor.



(b) Dining room.



(c) Kitchen.



(d) Conference room.

**LSGAN** [Mao et al., 2017]



**BEGAN** [Berthelot et al., 2017]

## Source → target domain transfer



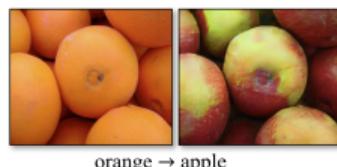
horse → zebra



zebra → horse



apple → orange



orange → apple

**CycleGAN** [Zhu et al., 2017]

## Text → image synthesis

this small bird has a pink breast and crown, and black primaries and secondaries.



this magnificent fellow is almost all black with a red crest, and white cheek patch.



[Zhu et al., 2017]

## Many GAN applications

### BW to Color



input

output

Labels to Street Scene



input

output

**Pix2Pix** [Isola et al., 2017]

# 2019: BigGAN



[Images source: Brock et al., 2019]

# “The GAN Zoo”

- GAN - Generative Adversarial Networks
- 3D-GAN - Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling
- acGAN - Face Aging With Conditional Generative Adversarial Networks
- AC-GAN - Conditional Image Synthesis With Auxiliary Classifier GANs
- AdaGAN - AdaGAN: Boosting Generative Models
- AEGAN - Learning Inverse Mapping by Autoencoder based Generative Adversarial Nets
- AffGAN - Amortised MAP Inference for Image Super-resolution
- AL-CGAN - Learning to Generate Images of Outdoor Scenes from Attributes and Semantic Layouts
- ALI - Adversarially Learned Inference
- AM-GAN - Generative Adversarial Nets with Labeled Data by Activation Maximization
- AnoGAN - Unsupervised Anomaly Detection with Generative Adversarial Networks to Guide Marker Discovery
- ArtGAN - ArtGAN: Artwork Synthesis with Conditional Categorical GANs
- b-GAN - b-GAN: Unified Framework of Generative Adversarial Networks
- Bayesian GAN - Deep and Hierarchical Implicit Models
- BEGAN - BEGAN: Boundary Equilibrium Generative Adversarial Networks
- BiGAN - Adversarial Feature Learning
- BS-GAN - Boundary-Seeking Generative Adversarial Networks
- CGAN - Conditional Generative Adversarial Nets
- CaloGAN - CaloGAN: Simulating 3D High Energy Particle Showers in Multi-Layer Electromagnetic Calorimeters with Generative Adversarial Networks
- CCGAN - Semi-Supervised Learning with Context-Conditional Generative Adversarial Networks
- CatGAN - Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks
- CoGAN - Coupled Generative Adversarial Networks
- Context-RNN-GAN - Contextual RNN-GANs for Abstract Reasoning Diagram Generation
- C-RNN-GAN - C-RNN-GAN: Continuous recurrent neural networks with adversarial training
- CS-GAN - Improving Neural Machine Translation with Conditional Sequence Generative Adversarial Nets
- CVAE-GAN - CVAE-GAN: Fine-Grained Image Generation through Asymmetric Training
- CycleGAN - Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks
- DTN - Unsupervised Cross-Domain Image Generation
- DCGAN - Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks
- DiscoGAN - Learning to Discover Cross-Domain Relations with Generative Adversarial Networks
- DR-GAN - Disentangled Representation Learning GAN for Pose-Invariant Face Recognition
- DualGAN - DualGAN: Unsupervised Dual Learning for Image-to-Image Translation
- EBGAN - Energy-based Generative Adversarial Network
- f-GAN - f-GAN: Training Generative Neural Samplers using Variational Divergence Minimization
- FF-GAN - Towards Large-Pose Face Frontalization in the Wild
- GAWWN - Learning What and Where to Draw
- GeneGAN - GeneGAN: Learning Object Transfiguration and Attribute Subspace from Unpaired Data
- Geometric GAN - Geometric GAN
- GoGAN - Gang of GANs: Generative Adversarial Networks with Maximum Margin Ranking
- GP-GAN - GP-GAN: Towards Realistic High-Resolution Image Blending
- IAN - Neural Photo Editing with Introspective Adversarial Networks
- iGAN - Generative Visual Manipulation on the Natural Image Manifold
- IcGAN - Invertible Conditional GANs for image editing
- ID-CGAN - Image De-raining Using a Conditional Generative Adversarial Network
- Improved GAN - Improved Techniques for Training GANs
- Infogan - InfoGAN: Interpretable Representation Learning by Information Maximizing Generative Adversarial Nets
- LAGAN - Learning Particle Physics by Example: Location-Aware Generative Adversarial Networks for Physics Synthesis
- LAPGAN - Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks

[Source: <https://github.com/hindupuravinash/the-gan-zoo>]

## Questions to Answer for Yourself / Discuss with Friends

- Repetition: While training a GAN, do we train the generator first or the discriminator first, and why not the other way around?
- Application of what you just learned: Can you think of other application domains, where GANs can be helpful, other than in computer vision?

# Lecture Overview

- 1 Introduction to Autoencoders
- 2 Sparse and Denoising Autoencoders
- 3 Generative Models and the Variational Autoencoder
- 4 Generative Adversarial Networks
- 5 Summary, Further Reading, References

Foundations of Deep Learning, Winter Term 2021/22

Week 12: Variational Autoencoders and Generative Adversarial Networks

Summary, Further Reading, References

Frank Hutter    Abhinav Valada

University of Freiburg



## Summary by Learning Goals

Having heard this lecture, you can now . . .

- explain the working mechanism behind autoencoders.
- describe regularization techniques for autoencoders such as sparsity and denoising.
- explain how to perform inference with variational autoencoders.
- elaborate on the reparametrization trick.
- describe the training mechanism of GANs.
- provide an overview of possible applications of GANs.

## Further Reading

- Resources for this lecture: Deep Learning Book [Goodfellow et al., 2016, chapter 14]

# References

- Berthelot, D., Schumm, T., Metz, Luke (2017)  
BEGAN: Boundary equilibrium generative adversarial networks  
*arXiv preprint arXiv:1703.10717*  
<https://arxiv.org/pdf/1703.10717.pdf>
- Brock, A., Donahue, J., Simonyan, K. (2019)  
Large Scale GAN Training for High Fidelity Natural Image Synthesis  
*International Conference on Learning Representations*  
<https://arxiv.org/pdf/1809.11096.pdf>
- Goodfellow, I., Bengio, Y., Courville, A. (2016)  
Deep Learning  
*MIT Press*  
<https://www.deeplearningbook.org>
- Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A., Bengio, Y. (2014)  
Generative adversarial nets  
*Advances in neural information processing systems*, 27  
<https://arxiv.org/pdf/1406.2661.pdf>
- Isola, P., Zhu, J.-Y., Zhou, T., Efros, A. (2017)  
Image-to-image translation with conditional adversarial networks  
*Proceedings of the IEEE conference on computer vision and pattern recognition*, 1125–1134  
<https://arxiv.org/pdf/1611.07004.pdf>
- Lange, S., Riedmiller, M., Voigtländer, A. (2012)  
Autonomous reinforcement learning on raw visual input data in a real world application  
*The 2012 international joint conference on neural networks (IJCNN)*, 1–8  
[https://ml.informatik.uni-freiburg.de/former/\\_media/publications/rieijcnn12.pdf](https://ml.informatik.uni-freiburg.de/former/_media/publications/rieijcnn12.pdf)
- Kingma, D., Welling, M. (2013)  
Auto-encoding variational bayes  
*arXiv preprint arXiv:1312.6114*  
<https://arxiv.org/pdf/1312.6114.pdf>
- Mao, X., Li, Q., Xie, H., Lau, R., Wang, Z., Paul Smolley, S. (2017)  
Least squares generative adversarial networks  
*Proceedings of the IEEE international conference on computer vision*, 2794–2802  
<https://arxiv.org/pdf/1611.04076.pdf>
- Radford, A., Metz, L., Chintala, S. (2016)  
Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks  
*4th International Conference on Learning Representations ICLR* (poster)  
<https://arxiv.org/pdf/1511.06434.pdf>
- Reed, S., Akata, Z., Yan, X., Logeswaran, L., Schiele, B., Lee, H. (2016)  
Generative adversarial text to image synthesis  
*International Conference on Machine Learning*, 1060–1069  
<https://arxiv.org/pdf/1605.05396.pdf>

# References

Zhu, J.-Y., Park, T., Isola, P., Efros, A. (2017)

Unpaired image-to-image translation using cycle-consistent adversarial networks

*Proceedings of the IEEE international conference on computer vision*, 2223–  
2232

<https://arxiv.org/pdf/1703.10593.pdf>