



Lecture 08: n-Step Bootstrapping and Eligibility Traces

75 plays · 105 players









 A public kahoot

Questions (6)

1 - Quiz

MC vs TD: What is correct?





60 sec

-  MC is more sensitive to initial values than TD. 
-  TD has a lower variance than MC. 
-  MC is usually more efficient than TD. 
-  MC has a larger bias. 

2 - Quiz

n-step returns: what is correct?







60 sec

-  $n=1$ is equivalent to the TD-target and $n=\infty$ to the MC-target. 
-  $n=1$ is equivalent to the MC-target and $n=\infty$ to the TD-target. 

3 - Quiz

Why are n-step targets not trivially applicable in Q-learning?

60 sec

-  The subtrajectory does not follow the target policy. 
-  n-step subreturn and current Q-estimation can have a different scale. 
-  The Q-prediction for some state s_{t+n} is less accurate than for s_{t+1} . 

4 - True or false

$\lambda=0.5$ is always the best choice for TD(λ).

30 sec



True



False



5 - True or false

In order to perform an update with λ -returns, we have to wait until the end of an episode.

30 sec



True



False



6 - Quiz

Eligibility Traces: what is correct?

60 sec



(1) encodes Frequency and (2) Recency.



(1) encodes Recency and (2) Frequency.

