

Lecture 8: Off-policy Methods with Function Approximation

Friday, December 17, 2021

Reinforcement Learning, Winter Term 2021/22

Joschka Boedecker, Gabriel Kalweit, Jasper Hoffmann

Neurorobotics Lab
University of Freiburg



Lecture Overview

- 1 Recap
- 2 Off-policy Learning with Function Approximation
- 3 Problems of Off-policy Learning with Function Approximation
- 4 Gradient-TD Methods
- 5 Deep Q-learning
- 6 Wrapup

Lecture Overview

- 1 Recap
- 2 Off-policy Learning with Function Approximation
- 3 Problems of Off-policy Learning with Function Approximation
- 4 Gradient-TD Methods
- 5 Deep Q-learning
- 6 Wrapup

Recap: Function Approximation in Reinforcement Learning

We want to update the weights w.r.t. the *Mean Squared Value Error* of the prediction:

$$\begin{aligned}\mathbf{w} &\leftarrow \mathbf{w} - \frac{1}{2}\alpha \nabla [v_\pi(S_t) - \hat{v}(S_t, \mathbf{w})]^2 \\ &\leftarrow \mathbf{w} + \alpha [v_\pi(S_t) - \hat{v}(S_t, \mathbf{w})] \nabla \hat{v}(S_t, \mathbf{w})\end{aligned}$$

However, we don't have $v_\pi(S_t)$.

Recap: Function Approximation in Reinforcement Learning

Gradient MC

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [G_t - \hat{v}(S_t, \mathbf{w})] \nabla \hat{v}(S_t, \mathbf{w})$$

Semi-gradient TD(0)

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w})] \nabla \hat{v}(S_t, \mathbf{w})$$

Why are bootstrapping methods, defined this way, called *semi-gradient methods*? They take into account the effects of changing \mathbf{w} w.r.t. the prediction, but not w.r.t. the target!

Recap: Semi-gradient SARSA

Episodic Semi-gradient Sarsa for Estimating $\hat{q} \approx q_*$

Input: a differentiable action-value function parameterization $\hat{q} : \mathcal{S} \times \mathcal{A} \times \mathbb{R}^d \rightarrow \mathbb{R}$

Algorithm parameters: step size $\alpha > 0$, small $\varepsilon > 0$

Initialize value-function weights $\mathbf{w} \in \mathbb{R}^d$ arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Loop for each episode:

$S, A \leftarrow$ initial state and action of episode (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'

 If S' is terminal:

$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R - \hat{q}(S, A, \mathbf{w})] \nabla \hat{q}(S, A, \mathbf{w})$

 Go to next episode

 Choose A' as a function of $\hat{q}(S', \cdot, \mathbf{w})$ (e.g., ε -greedy)

$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R + \gamma \hat{q}(S', A', \mathbf{w}) - \hat{q}(S, A, \mathbf{w})] \nabla \hat{q}(S, A, \mathbf{w})$

$S \leftarrow S'$

$A \leftarrow A'$

Lecture Overview

- 1 Recap
- 2 Off-policy Learning with Function Approximation**
- 3 Problems of Off-policy Learning with Function Approximation
- 4 Gradient-TD Methods
- 5 Deep Q-learning
- 6 Wrapup

- We want to learn the optimal policy, but we have to account for the problem of *maintaining exploration*
- We call the (optimal) policy to be learned the *target policy* π and the policy used to generate behaviour the *behaviour policy* b
- We say that learning is from data *off* the target policy – thus, those methods are referred to as *off-policy learning*
- Today: Off-policy learning methods with function approximation

Semi-gradient Off-policy TD(0)

Replace the update to an array to an update to weight vector \mathbf{w} .

Recap: Importance Sampling Ratio

$$\rho_t = \frac{\pi(A_t|S_t)}{b(A_t|S_t)}$$

Semi-gradient Off-policy TD(0)

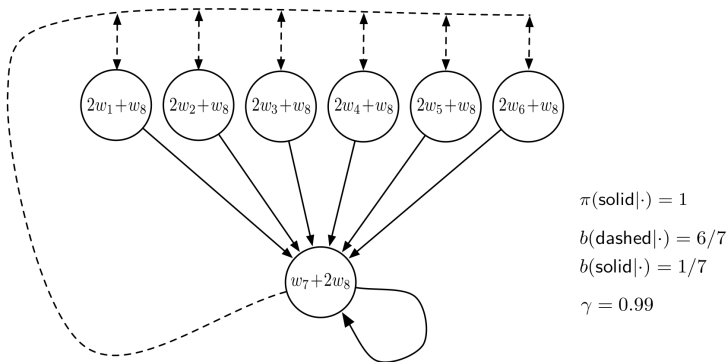
$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \rho_t \delta_t \nabla \hat{v}(S_t, \mathbf{w})$$

$$\text{where } \delta_t = R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w})$$

Lecture Overview

- 1 Recap
- 2 Off-policy Learning with Function Approximation
- 3 Problems of Off-policy Learning with Function Approximation**
- 4 Gradient-TD Methods
- 5 Deep Q-learning
- 6 Wrapup

Baird's Counterexample





The reward is 0 for all transitions, hence $v_\pi(s) = 0$. This could be exactly approximated by $\mathbf{w} = \mathbf{0}$.

Baird's Counterexample

Semi-gradient DP

$$\mathbf{w} \leftarrow \mathbf{w} + \frac{\alpha}{|\mathcal{S}|} \sum_{s \in \mathcal{S}} (\mathbb{E}[R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) | S_t = s] - \hat{v}(s, \mathbf{w})) \nabla \hat{v}(s, \mathbf{w})$$



The Deadly Triad

The combination of

- Function Approximation,
- Bootstrapping and
- Off-policy Learning

is known as the *Deadly Triad*, since it can lead to stability issues and divergence.

Lecture Overview

- 1 Recap
- 2 Off-policy Learning with Function Approximation
- 3 Problems of Off-policy Learning with Function Approximation
- 4 Gradient-TD Methods**
- 5 Deep Q-learning
- 6 Wrapup

Gradient-TD Methods

To this point, the TD-methods discussed did not leverage the true gradient (they are called semi-gradient methods). Recall the Bellman Equation:

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a)[r + \gamma v_{\pi}(s')]$$

Bellman Error

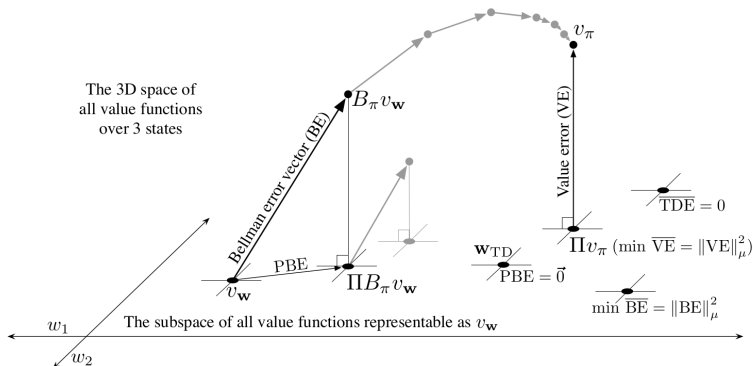
$$\begin{aligned} \bar{\delta}_{\mathbf{w}}(s) &= \left(\sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a)[r + \gamma v_{\mathbf{w}}(s')] \right) - v_{\mathbf{w}}(s) \\ &= \mathbb{E}_{\pi}[R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - v_{\mathbf{w}}(S_t) | S_t = s, A_t \sim \pi] \end{aligned}$$

Gradient-TD Methods

Projected Bellman Error

The Projected Bellman Error is the projection of the Bellman Error back into the representable space:

$$\overline{\text{PBE}} = \|\Pi \bar{\delta}_{\mathbf{w}}\|_{\mu}^2$$



Projection Matrix for linear FA

The projection matrix for linear FA can be represented as an $|\mathcal{S}| \times |\mathcal{S}|$ matrix:

$$\Pi = \mathbf{X}(\mathbf{X}^\top \mathbf{D} \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{D}$$

$$\begin{aligned} \overline{\text{PBE}}(\mathbf{w}) &= \|\Pi \bar{\delta}_{\mathbf{w}}\|_{\mu}^2 \\ &= (\Pi \bar{\delta}_{\mathbf{w}})^\top \mathbf{D} \Pi \bar{\delta}_{\mathbf{w}} \\ &= \bar{\delta}_{\mathbf{w}}^\top \Pi^\top \mathbf{D} \Pi \bar{\delta}_{\mathbf{w}} \\ &= \bar{\delta}_{\mathbf{w}}^\top \mathbf{D} \mathbf{X} (\mathbf{X}^\top \mathbf{D} \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{D} \bar{\delta}_{\mathbf{w}} \\ &= (\mathbf{X}^\top \mathbf{D} \bar{\delta}_{\mathbf{w}})^\top (\mathbf{X}^\top \mathbf{D} \mathbf{X})^{-1} (\mathbf{X}^\top \mathbf{D} \bar{\delta}_{\mathbf{w}}) \end{aligned}$$

The gradient with respect to \mathbf{w} is:

$$\nabla \overline{\text{PBE}}(\mathbf{w}) = 2 \nabla [\mathbf{X}^\top \mathbf{D} \bar{\delta}_{\mathbf{w}}]^\top (\mathbf{X}^\top \mathbf{D} \mathbf{X})^{-1} (\mathbf{X}^\top \mathbf{D} \bar{\delta}_{\mathbf{w}})$$

Gradient-TD Methods

Assume μ to be the distribution of state visited under the behaviour policy.

1

$$\mathbf{X}^\top \mathbf{D} \bar{\delta}_{\mathbf{w}} = \sum_s \mu(s) \mathbf{x}(s) \bar{\delta}_{\mathbf{w}}(s) = \mathbb{E}[\rho_t \delta_t \mathbf{x}_t]$$

2

$$\begin{aligned} \nabla[\mathbf{X}^\top \mathbf{D} \bar{\delta}_{\mathbf{w}}]^\top &= \mathbb{E}[\rho_t \nabla \delta_t^\top \mathbf{x}_t^\top] \\ &= \mathbb{E}[\rho_t \nabla (R_{t+1} + \gamma \mathbf{w}^\top \mathbf{x}_{t+1} - \mathbf{w}^\top \mathbf{x}_t)^\top \mathbf{x}_t^\top] \\ &= \mathbb{E}[\rho_t (\gamma \mathbf{x}_{t+1} - \mathbf{x}_t) \mathbf{x}_t^\top] \end{aligned}$$

3


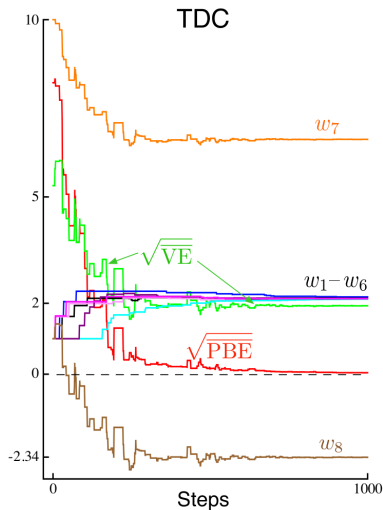
$$\mathbf{X}^\top \mathbf{D} \mathbf{X} = \sum_s \mu(s) \mathbf{x}_s \mathbf{x}_s^\top = \mathbb{E}[\mathbf{x}_t \mathbf{x}_t^\top]$$

- We can thus rewrite the gradient as:

$$\nabla \overline{\text{PBE}}(\mathbf{w}) = 2\mathbb{E}[\rho_t(\gamma\mathbf{x}_{t+1} - \mathbf{x}_t)\mathbf{x}_t^\top]\mathbb{E}[\mathbf{x}_t\mathbf{x}_t^\top]\mathbb{E}[\rho_t\delta_t\mathbf{x}_t]$$

- Store the last two factors in d -vector $\mathbf{v} \approx \mathbb{E}[\mathbf{x}_t\mathbf{x}_t^\top]^{-1}\mathbb{E}[\rho_t\delta_t\mathbf{x}_t]$ and update via $\mathbf{v} \leftarrow \mathbf{v} + \beta\rho_t(\delta_t - \mathbf{v}^\top\mathbf{x}_t)\mathbf{x}_t$
- Estimate the first part via sampling
- This method and variants of it are called Gradient-TD methods

Gradient-TD Methods



Lecture Overview

- 1 Recap
- 2 Off-policy Learning with Function Approximation
- 3 Problems of Off-policy Learning with Function Approximation
- 4 Gradient-TD Methods
- 5 Deep Q-learning**
- 6 Wrapup

Neural Fitted-Q Iteration (NFQ) [Riedmiller 2005]

- Model-free off-policy RL algorithm that works on continuous state and discrete action spaces
- Q-function is represented by a multi-layer perceptron
- One of the first approaches that combined RL with ANNs, predecessor of DQN

Neural Fitted-Q Iteration (NFQ) [Riedmiller 2005]

Algorithm 1 NFQ

for iteration $i = 1, \dots, N$ **do**

sample trajectory with ϵ -greedy exploration and add to memory D

initialize network weights randomly

generate pattern set $P = \{(x_j, y_j) | j = 1..|D|\}$ with

$x_j = (s_j, a_j)$ and $y_j = \begin{cases} r_j & \text{if } s_j \text{ is terminal} \\ r_j + \gamma \max_{a'} Q(s_{j+1}, a', \mathbf{w}_i) & \text{else} \end{cases}$

for iteration $k = 1, \dots, K$ **do**

Fit weights according to:

$$L(\mathbf{w}_i) = \frac{1}{|D|} \sum_{j=1}^{|D|} (y_j - Q(x_j, \mathbf{w}_i))^2$$

Deep Q-Networks (DQN)

DQN provides a stable solution to deep RL:

- Use experience replay (as in NFQ)
- Sample minibatches (as opposed to Full Batch in NFQ)
- Freeze target Q-networks (no target networks in NFQ)
- Optional: Clip rewards or normalize network adaptively to sensible range

Deep Q-Networks: Experience Replay

To remove correlations, build data set from agent's own experience

- Take action a_t according to ϵ -greedy policy
- Store transition $(s_t, a_t, r_{t+1}, s_{t+1})$ in replay memory D
- Sample random mini-batch of transitions (s, a, r, s') from D
- Optimize MSE between Q-network and Q-learning targets, e.g.

$$L(\mathbf{w}) = \mathbb{E}_{s,a,r,s' \sim D} [(r + \gamma \max_{a'} Q(s', a', \mathbf{w}) - Q(s, a, \mathbf{w}))^2]$$

Deep Q-Networks: Target Networks

To avoid oscillations, fix parameters used in Q-learning target

- Compute Q-learning targets w.r.t. old, fixed parameters \mathbf{w}^-

$$r + \gamma \arg \max_{a'} Q(s', a', \mathbf{w}^-)$$

- Optimize MSE between Q-network and Q-learning targets

$$L(\mathbf{w}) = \mathbb{E}_{s,a,r,s' \sim D} [(r + \gamma \max_{a'} Q(s', a', \mathbf{w}^-) - Q(s, a, \mathbf{w}))^2]$$

- Periodically update fixed parameters $\mathbf{w}^- \leftarrow \mathbf{w}$
 - hard update: update target network every N steps
 - slow update: slowly update weights of target network every step by

$$\mathbf{w}^- \leftarrow (1 - \tau)\mathbf{w}^- + \tau\mathbf{w}$$

Deep Q-Networks (DQN)

Algorithm 2 DQN

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights

for episode $i = 1, \dots, M$ **do**

for $t = 1, \dots, T$ **do**

 select action a_t ϵ -greedily

 Store transition (s_t, a_t, s_{t+1}, r_t) in D

 Sample minibatch of transitions (s_j, a_j, r_j, s_{j+1}) from D

 Set $y_j = \begin{cases} r_j & \text{if } s_{j+1} \text{ is terminal} \\ r_j + \gamma \max_{a'} Q(s_{j+1}, a', \mathbf{w}^-) & \text{else} \end{cases}$


 Update the parameters of Q according to:

$$\begin{aligned} \nabla_{\mathbf{w}_i} L_i(\mathbf{w}_i) = & \mathbb{E}_{s,a,s',r \sim D} [(r + \gamma \max_{a'} Q(s', a', \mathbf{w}_i) \\ & - Q(s, a, \mathbf{w}_i)) \nabla_{\mathbf{w}_i} Q(s, a, \mathbf{w}_i)] \end{aligned}$$

 Update target network

This is your
exercise

Deep Q-Networks: Reinforcement Learning in Atari



Deep Q-Networks: Reinforcement Learning in Atari

- End-to-end learning of values $Q(s, a)$ from pixels s
- Input state s is a stack of raw pixels from the last 4 frames
- Output is $Q(s, a)$ for 18 joystick/button positions
- Reward is change in score for that step

How much does DQN help?

	Q-Learning	Q-Learning + Target Q	Q-Learning + Replay	DQN Q-learning + Replay + Target Q
Breakout	3	10	241	317
Enduro	29	142	831	1006
River Raid	1453	2868	4103	7447
Seaquest	276	1003	831	2894
Space Invaders	302	373	826	1089

Lecture Overview

- 1 Recap
- 2 Off-policy Learning with Function Approximation
- 3 Problems of Off-policy Learning with Function Approximation
- 4 Gradient-TD Methods
- 5 Deep Q-learning
- 6 Wrapup**

Summary by Learning Goals

Having heard this lecture, you can now...

- Explain the difficulties that may arise in off-policy learning with FA
- Apply deep Q-learning and explain why we include replay memory and target networks