

Lecture 9: Eligibility Traces

Friday, January 14, 2022

Reinforcement Learning, Winter Term 2021/22

Joschka Boedecker, Gabriel Kalweit, Jasper Hoffmann

Neurorobotics Lab
University of Freiburg



Lecture Overview

- 1 Recap
- 2 n -step Bootstrapping
- 3 Eligibility Traces
- 4 Wrapup

Lecture Overview

- 1 Recap
- 2 n -step Bootstrapping
- 3 Eligibility Traces
- 4 Wrapup

Recap: Off-policy Methods with Function Approximation

- *Deadly Triad*: The combination of Function Approximation, Bootstrapping and Off-policy Learning
- Gradient TD methods: TD(0) with gradient correction (TDC)
- NFQ (experience replay, full batch)
- DQN (experience replay, minibatches, target networks)

Lecture Overview

- 1 Recap
- 2 n -step Bootstrapping
- 3 Eligibility Traces
- 4 Wrapup

n -step Bootstrapping

- Recall:

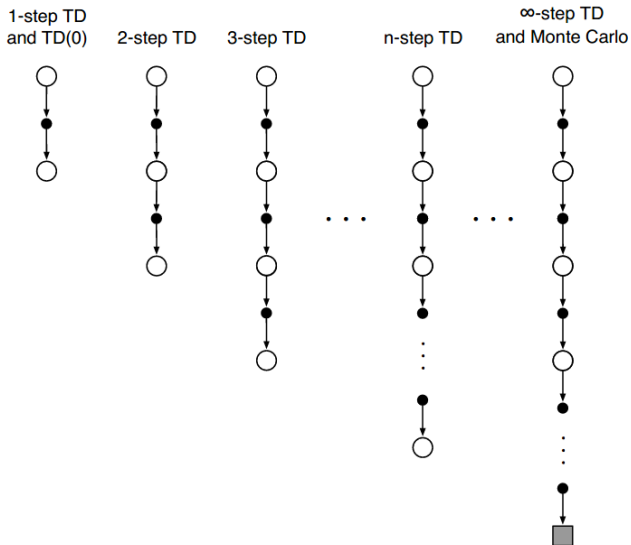
- MC-target: $R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-1} R_T$
- TD-target: $R_{t+1} + \gamma V(S_{t+1})$

- n -step Return:

$$G_{t:t+n} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$$

- $n = 0$: TD
- $n = \infty$: MC

n -step TD Prediction



n -step TD Prediction

n -step TD for estimating $V \approx v_\pi$

Input: a policy π

Algorithm parameters: step size $\alpha \in (0, 1]$, a positive integer n

Initialize $V(s)$ arbitrarily, for all $s \in \mathcal{S}$

All store and access operations (for S_t and R_t) can take their index mod $n + 1$

Loop for each episode:

 Initialize and store $S_0 \neq \text{terminal}$

$T \leftarrow \infty$

 Loop for $t = 0, 1, 2, \dots$:

 If $t < T$, then:

 Take an action according to $\pi(\cdot | S_t)$

 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

 If S_{t+1} is terminal, then $T \leftarrow t + 1$

$\tau \leftarrow t - n + 1$ (τ is the time whose state's estimate is being updated)

 If $\tau \geq 0$:

$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$

 If $\tau + n < T$, then: $G \leftarrow G + \gamma^n V(S_{\tau+n})$ ($G_{\tau:\tau+n}$)

$V(S_\tau) \leftarrow V(S_\tau) + \alpha [G - V(S_\tau)]$

 Until $\tau = T - 1$

How can n -step methods be used not just for prediction but for control?

$$G_{t:t+n} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q_{t+n-1}(S_{t+n}, A_{t+n})$$

where $n \geq 1$ and $0 \leq t < T - n$ and $G_{t:t+n} = G_t$ if $t + n \geq T$.

n -step Sarsa

$$Q_{t+n}(S_t, A_t) = Q_{t+n-1}(S_t, A_t) + \alpha[G_{t:t+n} - Q_{t+n-1}(S_t, A_t)], \quad 0 \leq t < T$$

while the values of all other states remain unchanged:

$$Q_{t+n}(s, a) = Q_{t+n-1}(s, a) \quad \forall s, a \text{ s.t. } s \neq S_t \text{ or } a \neq A_t$$

n -step Sarsa

1-step Sarsa
aka Sarsa(0)



2-step Sarsa



3-step Sarsa



...

n -step Sarsa



∞ -step Sarsa
aka Monte Carlo




n -step
Expected Sarsa




Example: Gridworld


Path taken



Action values increased
by one-step Sarsa



Action values increased
by 10-step Sarsa



n -step Off-policy Learning

Importance sampling ratio (the relative probability under the two policies of taking the n actions from A_t to A_{t+n-1}):

$$\rho_{t:h} = \prod_{k=t}^{\min(h, T-1)} \frac{\pi(A_k|S_k)}{b(A_k|S_k)}$$

Previous n -step Sarsa update can be replaced by an off-policy form:

$$Q_{t+n}(S_t, A_t) = Q_{t+n-1}(S_t, A_t) + \alpha \rho_{t+1:t+n} [G_{t:t+n} - Q_{t+n-1}(S_t, A_t)]$$

with $0 \leq t < T$.

Off-policy n -step Sarsa

Off-policy n -step Sarsa for estimating $Q \approx q_*$ or q_π

Input: an arbitrary behavior policy b such that $b(a|s) > 0$, for all $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize $Q(s, a)$ arbitrarily, for all $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize π to be greedy with respect to Q , or as a fixed given policy

Algorithm parameters: step size $\alpha \in (0, 1]$, a positive integer n

All store and access operations (for S_t , A_t , and R_t) can take their index mod $n + 1$

Loop for each episode:

 Initialize and store $S_0 \neq \text{terminal}$

 Select and store an action $A_0 \sim b(\cdot|S_0)$

$T \leftarrow \infty$

 Loop for $t = 0, 1, 2, \dots$:

 If $t < T$, then:

 Take action A_t

 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

 If S_{t+1} is terminal, then:

$T \leftarrow t + 1$

 else:

 Select and store an action $A_{t+1} \sim b(\cdot|S_{t+1})$

$\tau \leftarrow t - n + 1$ (τ is the time whose estimate is being updated)

 If $\tau \geq 0$:

$\rho \leftarrow \prod_{i=\tau+1}^{\min(\tau+n-1, T-1)} \frac{\pi(A_i|S_i)}{b(A_i|S_i)}$ ($\rho_{\tau+1:t+n-1}$)

$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$ ($G_{\tau:\tau+n}$)

 If $\tau + n < T$, then: $G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n})$ ($G_{\tau:\tau+n}$)

$Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha \rho [G - Q(S_\tau, A_\tau)]$

 If π is being learned, then ensure that $\pi(\cdot|S_\tau)$ is greedy wrt Q

 Until $\tau = T - 1$

Averaging n -step Returns

- We can average n -step returns over different n
- e.g. average the 2-step and 4-step returns

$$\frac{1}{2}G_{t:t+2} + \frac{1}{2}G_{t:t+4}$$

- Combines information from two different time steps
- Can we efficiently combine information from all time-steps?

Lecture Overview

- 1 Recap
- 2 n -step Bootstrapping
- 3 Eligibility Traces**
- 4 Wrapup

Eligibility Traces and λ -return

Eligibility traces unify and generalize TD and Monte Carlo methods


- MC methods at one end ($\lambda = 1$) and one-step TD methods at the other ($\lambda = 0$)
- almost any temporal-difference (TD) method can be combined with eligibility traces to (maybe) learn more efficiently

λ -return

- For infinite control tasks: $G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_{t:t+n}$
- For episodic control tasks:

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{T-t-1} \lambda^{n-1} G_{t:t+n} + \lambda^{T-t-1} G_t$$

λ -return



Offline λ -Return Algorithm


At the end of the episode, updates are made according to:

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t^\lambda - V(S_t)),$$

or, in the FA setting, to the semi-gradient rule:

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha[G_t^\lambda - \hat{v}(S_t, \mathbf{w}_t)]\nabla\hat{v}(S_t, \mathbf{w}_t), t = 0, \dots, T - 1$$

Alternative way of moving smoothly between MC and one-step TD, can be compared to n -step TD methods.



Offline λ -Return Algorithm

The Offline λ -update can be converted to TD-form:

$$V(S_t) \leftarrow V(S_t) + \alpha \sum_{i=0}^{\infty} \lambda^i \delta_{t+i+1}$$

Lemma 9.1

$$(1 - \lambda) \sum_{n=m}^{\infty} \lambda^{n-1} = \lambda^{m-1}$$

Proof.

$$\begin{aligned} (1 - \lambda) \sum_{n=m}^{\infty} \lambda^{n-1} &= (1 - \lambda) \left(\sum_{n=1}^{\infty} \lambda^{n-1} - \sum_{k=1}^{m-1} \lambda^{k-1} \right) \\ &= (1 - \lambda) \left(\sum_{n=0}^{\infty} \lambda^n - \sum_{k=0}^{m-2} \lambda^k \right) \\ &= (1 - \lambda) \left(\frac{1}{1 - \lambda} - \frac{1 - \lambda^{m-1}}{1 - \lambda} \right) \\ &= 1 - (1 - \lambda^{m-1}) \\ &= \lambda^{m-1} \end{aligned}$$

Offline λ -Return Algorithm

Lemma 9.2

$$\sum_{n=1}^{\infty} \lambda^{n-1} \sum_{m=1}^n R_{t+m} = \sum_{m=1}^{\infty} R_{t+m} \sum_{n=m}^{\infty} \lambda^{n-1}$$

Proof.

$$\begin{aligned} \sum_{n=1}^{\infty} \lambda^{n-1} \sum_{m=1}^n R_{t+m} &= \lambda^0 R_{t+1} + \lambda^1 (R_{t+1} + R_{t+2}) \\ &\quad + \lambda^2 (R_{t+1} + R_{t+2} + R_{t+3}) + \dots \\ &= R_{t+1} (\lambda^0 + \lambda^1 + \lambda^2 + \dots) + R_{t+2} (\lambda^1 + \lambda^2 + \dots) \\ &\quad + R_{t+3} (\lambda^2 + \dots) + \dots \\ &= \sum_{m=1}^{\infty} R_{t+m} \sum_{n=m}^{\infty} \lambda^{n-1} \end{aligned}$$

□

Offline λ -Return Algorithm

Lemma 9.3

$$(1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} \sum_{m=1}^n R_{t+m} = \sum_{m=1}^{\infty} \lambda^{m-1} R_{t+m}$$

Proof.

$$\begin{aligned} (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} \sum_{m=1}^n R_{t+m} &\stackrel{2}{=} (1 - \lambda) \sum_{m=1}^{\infty} R_{t+m} \sum_{n=m}^{\infty} \lambda^{n-1} \\ &= \sum_{m=1}^{\infty} R_{t+m} (1 - \lambda) \sum_{n=m}^{\infty} \lambda^{n-1} \\ &\stackrel{1}{=} \sum_{m=1}^{\infty} \lambda^{m-1} R_{t+m} \end{aligned}$$

□

Offline λ -Return Algorithm

Lemma 9.4

$$(1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} V(S_{t+n}) = V(S_t) + \sum_{m=1}^{\infty} \lambda^{m-1} [V(S_{t+m}) - V(S_{t+m-1})]$$

Proof.

$$\begin{aligned} (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} V(S_{t+n}) &= \sum_{n=1}^{\infty} V(S_{t+n}) (\lambda^{n-1} - \lambda^n) \\ &= V(S_{t+1})(\lambda^0 - \lambda^1) + V(S_{t+2})(\lambda^1 - \lambda^2) + \dots \\ &= \lambda^0 V(S_{t+1}) + \lambda^1 V(S_{t+2}) - \lambda^1 V(S_{t+1}) + \\ &\quad \lambda^2 V(S_{t+3}) - \lambda^2 V(S_{t+2}) + \dots + \\ &\quad \lambda^0 V(S_t) - \lambda^0 V(S_t) \\ &= V(S_t) + \sum_{m=1}^{\infty} \lambda^{m-1} [V(S_{t+m}) - V(S_{t+m-1})] \end{aligned}$$

Offline λ -Return Algorithm

The Offline λ -update can be converted to TD-form:

$$V(S_t) \leftarrow V(S_t) + \alpha \sum_{i=0}^{\infty} \lambda^i \delta_{t+i+1}$$

Proof. For simplification, we will ignore the discount.

$$\begin{aligned} V(S_t) &= (1 - \lambda) \cdot \mathbb{E} \left[\sum_{n=1}^{\infty} \lambda^{n-1} G_{t:t+n} \right] \\ &= (1 - \lambda) \cdot \mathbb{E} \left[\sum_{n=1}^{\infty} \lambda^{n-1} \left(\sum_{m=1}^n R_{t+m} + V(S_{t+n}) \right) \right] \\ &\stackrel{3,4}{=} \mathbb{E} \left[\sum_{m=1}^{\infty} \lambda^{m-1} [R_{t+m} + V(S_{t+m}) - V(S_{t+m-1})] + V(S_t) \right] \\ &= \mathbb{E} \left[\sum_{m=1}^{\infty} \lambda^{m-1} \delta_{t+m} + V(S_t) \right] \end{aligned}$$

Offline λ -Return Algorithm

The Offline λ -update can be converted to TD-form:

$$V(S_t) \leftarrow V(S_t) + \alpha \sum_{i=0}^{\infty} \lambda^i \delta_{t+i+1}$$

Proof. For simplification, we will ignore the discount.
If we plug this into our value function update, we get:


$$V(S_t) \leftarrow V(S_t) + \alpha \left(\sum_{m=1}^{\infty} \lambda^{m-1} \delta_{t+m} + V(S_t) - V(S_t) \right),$$

which leads to:

$$V(S_t) \leftarrow V(S_t) + \alpha \sum_{m=0}^{\infty} \lambda^m \delta_{t+m+1}.$$




Forward View



- Update value function towards the λ -return
- Forward-view looks into the future to compute G_t^λ
- Like MC, can only be computed from complete episodes

Backward View



- look at the current TD error δ_t
- assign it backward to each prior state according to how much that state contributed to the current eligibility trace at that time

Eligibility Traces

- Eligibility Traces assign credit to components of the weight vector according to their contribution to state valuations
- They combine heuristics of *Frequency* and *Recency* (implemented by a λ -decay)
- With function approximation, the eligibility trace is a vector $\mathbf{z}_t \in \mathbb{R}$, initialized by $\mathbf{z}_{-1} = \mathbf{0}$ and incremented on each time step by:

$$\mathbf{z}_t = \gamma \lambda \mathbf{z}_{t-1} + \nabla \hat{v}(S_t, \mathbf{w}_t), \quad 0 \leq t \leq T,$$

where λ is called trace-decay parameter.

- TD(λ) updates the weight vector on every step of an episode rather than only at the end and is not limited to episodic problems
- With the TD error

$$\delta_t = R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}_t) - \hat{v}(S_t, \mathbf{w}_t),$$

in TD(λ), the weight vector is updated by:

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha \delta_t \mathbf{z}_t.$$

Semi-gradient TD(λ) for estimating $\hat{v} \approx v_\pi$

Input: the policy π to be evaluated

Input: a differentiable function $\hat{v} : \mathcal{S}^+ \times \mathbb{R}^d \rightarrow \mathbb{R}$ such that $\hat{v}(\text{terminal}, \cdot) = 0$

Algorithm parameters: step size $\alpha > 0$, trace decay rate $\lambda \in [0, 1]$

Initialize value-function weights \mathbf{w} arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Loop for each episode:

 Initialize S

$\mathbf{z} \leftarrow \mathbf{0}$

(a d -dimensional vector)

 Loop for each step of episode:

 | Choose $A \sim \pi(\cdot|S)$

 | Take action A , observe R, S'

 | $\mathbf{z} \leftarrow \gamma\lambda\mathbf{z} + \nabla\hat{v}(S, \mathbf{w})$

 | $\delta \leftarrow R + \gamma\hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})$

 | $\mathbf{w} \leftarrow \mathbf{w} + \alpha\delta\mathbf{z}$

 | $S \leftarrow S'$

 until S' is terminal

n -step Truncated λ -return

- offline λ -return algorithm is limited because:
 - the λ -return is not known until the end of the episode, or
 - technically never known in the continuing case
- dependence becomes weaker for longer-delayed reward, falling by $\gamma\lambda$ for each step of delay
- natural approximation: truncate the sequence after some number of steps:

truncated λ -return

$$G_{t:h}^{\lambda} = (1 - \lambda) \sum_{n=1}^{h-t-1} \lambda^{n-1} G_{t:t+n} + \lambda^{h-t-1} G_{t:h}, \quad 0 \leq t < h \leq T$$

n -step Truncated λ -return

- Now: updates are delayed by n steps and only take into account the first n rewards, but now all the k -step returns are included for $1 \leq k \leq n$.
- State-value case: truncated $\text{TD}(\lambda)$ or $\text{TTD}(\lambda)$

$\text{TTD}(\lambda)$

$$\mathbf{w}_{t+n} = \mathbf{w}_{t+n-1} + \alpha [G_{t:t+n}^\lambda - \hat{v}(S_t, \mathbf{w}_{t+n-1})] \nabla \hat{v}(S_t, \mathbf{w}_{t+n-1}), 0 \leq t < T.$$

The k -step λ -return can efficiently implemented as:

$$G_{t:t+k} = \hat{v}(S_t, \mathbf{w}_{t-1}) + \sum_{i=t}^{t+k-1} (\gamma\lambda)^{i-t} \delta'_i$$

with $\delta'_i = R_{t+1} + \gamma\hat{v}(S_{t+1}, \mathbf{w}_{t-1}) - \hat{v}(S_t, \mathbf{w}_{t-1})$.

Extension of eligibility traces to action-value methods.

- Action-value form of the n -step return:

$$G_{t:t+n} = R_{t+1} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n \hat{q}(S_{t+n}, A_{t+n}, \mathbf{w}_{t+n-1}), \quad t+n < T$$

with $G_{t:t+n} = G_t$ if $t+n \geq T$.

- Action-value form of the truncated λ -return:

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [G_t^\lambda - \hat{q}(S_t, A_t, \mathbf{w}_t)] \nabla \hat{q}(S_t, A_t, \mathbf{w}_t), \quad t = 0, \dots, T-1,$$

where $G_t^\lambda = G_{t:\infty}^\lambda$.

- With the TD error

$$\delta_t = R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}, \mathbf{w}_t) - \hat{q}(S_t, A_t, \mathbf{w}_t),$$

- And the action-value form of the TD error:

$$\mathbf{z}_{-1} = \mathbf{0}$$


$$\mathbf{z}_t = \gamma \lambda \mathbf{z}_{t-1} + \nabla \hat{q}(S_t, A_t, \mathbf{w}_t), \quad 0 \leq t \leq T,$$

in Sarsa(λ), the weight vector is updated by:


$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha \delta_t \mathbf{z}_t.$$

Example: Traces in Gridworld


Path taken




Action values increased by one-step Sarsa



Action values increased
by 10-step Sarsa



Action values increased by Sarsa(λ) with $\lambda=0.9$



Lecture Overview

- 1 Recap
- 2 n -step Bootstrapping
- 3 Eligibility Traces
- 4 Wrapup**

Summary by Learning Goals

Having heard this lecture, you can now...

- explain two different ways of shifting and choosing between Monte Carlo and TD methods
- why eligibility trace methods are more general and often faster to learn